

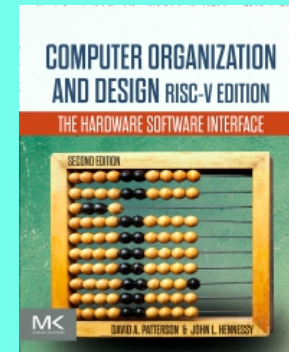
CSC 411

Computer Organization (Spring 2022)
Lecture 5b: RISC-V Memory Organization

Prof. Marco Alvarez, University of Rhode Island

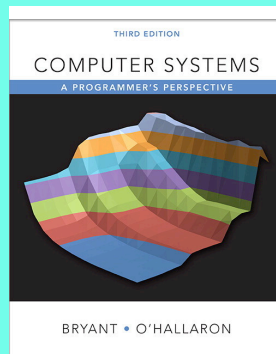
Disclaimer

Some of the following slides are adapted from:
Computer Organization and Design (Patterson and Hennessy)
The Hardware/Software Interface



Disclaimer

Some of the following slides are copied from:
Computer Systems (Bryant and O'Hallaron)
A Programmer's Perspective



Byte-Oriented Memory Organization



■ Programs refer to data by address

- Conceptually, envision it as a very large array of bytes
 - In reality, it's not, but can think of it that way
- An address is like an index into that array
 - and, a pointer variable stores an address

■ Note: system provides private address spaces to each "process"

- Think of a process as a program being executed
- So, a program can clobber its own data, but not that of others

Machine Words

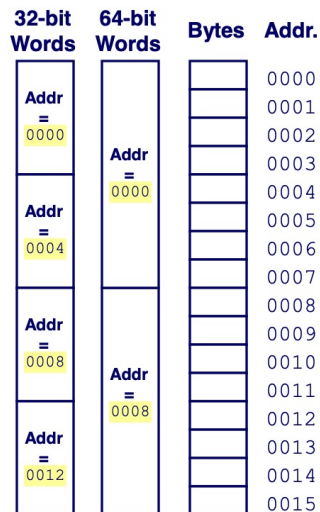
- **Any given computer has a “Word Size”**
 - Nominal size of integer-valued data
 - and of addresses
 - Until recently, most machines used 32 bits (4 bytes) as word size
 - Limits addresses to 4GB (2^{32} bytes)
 - Increasingly, machines have 64-bit word size
 - Potentially, could have 18 EB (exabytes) of addressable memory
 - That's 18.4×10^{18}
 - Machines still support multiple data formats
 - Fractions or multiples of word size
 - Always integral number of bytes

Example Data Representations

C Data Type	Typical 32-bit	Typical 64-bit	x86-64
char	1	1	1
short	2	2	2
int	4	4	4
long	4	8	8
float	4	4	4
double	8	8	8
pointer	4	8	8

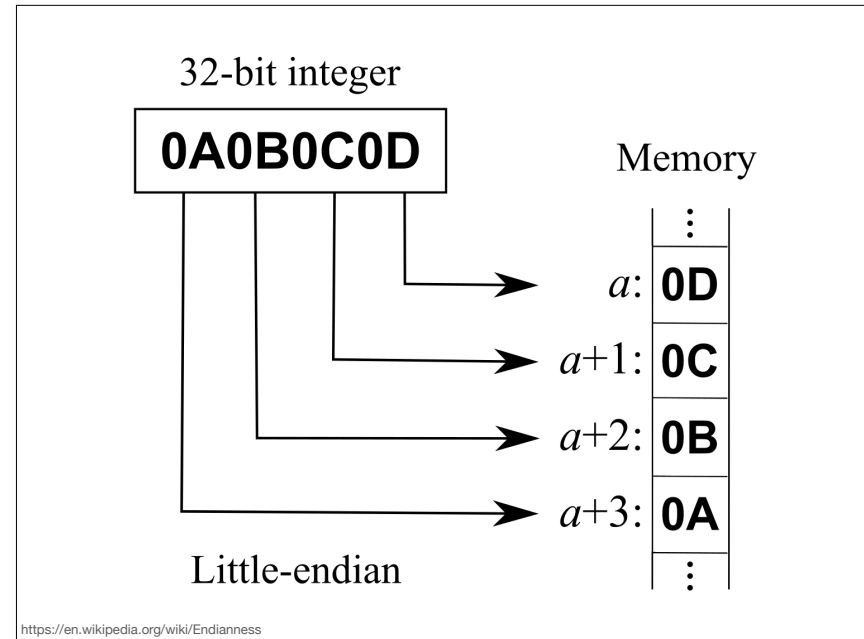
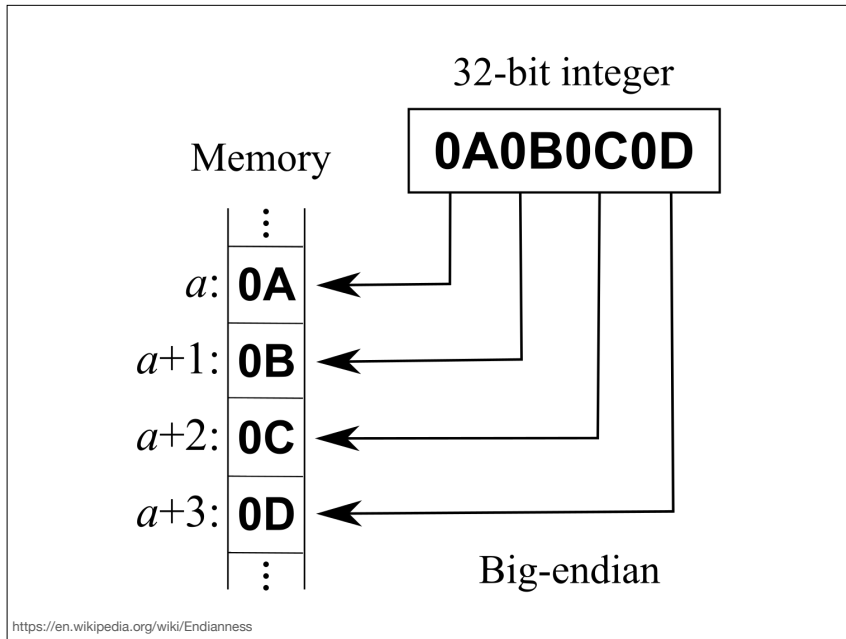
Word-Oriented Memory Organization

- **Addresses Specify Byte Locations**
 - Address of first byte in word
 - Addresses of successive words differ by 4 (32-bit) or 8 (64-bit)



Byte Ordering

- **So, how are the bytes within a multi-byte word ordered in memory?**
- **Conventions**
 - Big Endian: Sun (Oracle SPARC), PPC Mac, *Internet*
 - Least significant byte has highest address
 - Little Endian: *x86*, ARM processors running Android, iOS, and Linux
 - Least significant byte has lowest address



Carnegie Mellon

Byte Ordering Example

- Example**
 - Variable x has 4-byte value of 0x01234567
 - Address given by &x is 0x100

Big Endian

0x100	0x101	0x102	0x103
01	23	45	67

Little Endian

0x100	0x101	0x102	0x103
67	45	23	01

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

42

Carnegie Mellon

Representing Integers

Decimal: 15213
Binary: 0011 1011 0110 1101
Hex: 3 B 6 D

int A = 15213;

long int C = 15213;

int B = -15213;

Two's complement representation

Increasing addresses ↓

IA32, x86-64 Sun

IA32 x86-64 Sun

Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

43

Examining Data Representations

Code to Print Byte Representation of Data

- Casting pointer to unsigned char * allows treatment as a byte array

```
typedef unsigned char *pointer;

void show_bytes(pointer start, size_t len){
    size_t i;
    for (i = 0; i < len; i++){
        printf("%p\t0x%.2x\n", start+i, start[i]);
        printf("\n");
    }
}
```

Printf directives:

%p: Print pointer
%x: Print Hexadecimal

show_bytes Execution Example

```
int a = 15213;
printf("int a = 15213;\n");
show_bytes((pointer) &a, sizeof(int));
```

Result (Linux x86-64):

```
int a = 15213;
0x7fffb7f71dbc 6d
0x7fffb7f71dbd 3b
0x7fffb7f71dbe 00
0x7fffb7f71dbf 00
```

Representing Strings

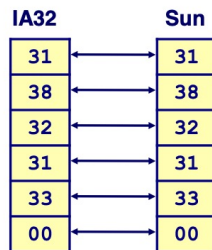
Strings in C

- Represented by array of characters
- Each character encoded in ASCII format
 - Standard 7-bit encoding of character set
 - Character "0" has code 0x30
 - Digit i has code $0x30+i$
- String should be null-terminated
 - Final character = 0

Compatibility

- Byte ordering not an issue

```
char S[6] = "18213";
```



ASCII control characters				ASCII printable characters				Extended ASCII characters			
00	NULL	(Null character)		32	space	64	@	96	.		
01	SOH	(Start of Header)		33	!	65	A	97	a		
02	STX	(Start of Text)		34	"	66	B	98	b		
03	ETX	(End of Text)		35	#	67	C	99	c		
04	EOT	(End of Trans.)		36	\$	68	D	100	d		
05	ENQ	(Enquiry)		37	%	69	E	101	e		
06	ACK	(Acknowledgement)		38	&	70	F	102	f		
07	BEL	(Bell)		39	'	71	G	103	g		
08	BS	(Backspace)		40	(72	H	104	h		
09	HT	(Horizontal Tab)		41)	73	I	105	i		
10	LF	(Line feed)		42	*	74	J	106	j		
11	VT	(Vertical Tab)		43	+	75	K	107	k		
12	FF	(Form feed)		44	,	76	L	108	l		
13	CR	(Carriage return)		45	-	77	M	109	m		
14	SO	(Shift Out)		46	.	78	N	110	n		
15	SI	(Shift In)		47	/	79	O	111	o		
16	DLE	(Data link escape)		48	0	80	P	112	p		
17	DC1	(Device control 1)		49	1	81	Q	113	q		
18	DC2	(Device control 2)		50	2	82	R	114	r		
19	DC3	(Device control 3)		51	3	83	S	115	s		
20	DC4	(Device control 4)		52	4	84	T	116	t		
21	NAK	(Negative acknowl.)		53	5	85	U	117	u		
22	SYN	(Synchronous idle)		54	6	86	V	118	v		
23	ETB	(End of trans. block)		55	7	87	W	119	w		
24	CAN	(Cancel)		56	8	88	X	120	x		
25	EM	(End of medium)		57	9	89	Y	121	y		
26	SUB	(Substitute)		58	:	90	Z	122	z		
27	ESC	(Escape)		59	;	91	[123	{		
28	FS	(File separator)		60	<	92	\	124			
29	GS	(Group separator)		61	=	93]	125	}		
30	RS	(Record separator)		62	>	94	^	126	~		
31	US	(Unit separator)		63	?	95	_				
127	DEL	(Delete)									
128	Ç			160	à			192	À		
129	ù			161	á			193	Á		
130	é			162	â			194	Â		
131	â			163	ã			195	Ã		
132	ä			164	ä			196	Ä		
133	å			165	Å			197	Å		
134	ä			166	ä			198	ä		
135	ç			167	ç			199	Ç		
136	è			168	è			200	È		
137	é			169	é			201	É		
138	ê			170	ê			202	Ê		
139	ï			171	ï			203	Ï		
140	ì			172	ì			204	Ì		
141	í			173	í			205	Í		
142	Ä			174	«			206	»		
143	Å			175	»			207	»		
144	É			176	»			208	»		
145	æ			177	»			209	»		
146	Æ			178	»			210	»		
147	ø			179	»			211	»		
148	ø			180	»			212	»		
149	ø			181	»			213	»		
150	ù			182	»			214	»		
151	ù			183	»			215	»		
152	ý			184	»			216	»		
153	Û			185	»			217	»		
154	Ü			186	»			218	»		
155	ü			187	»			219	»		
156	ø			188	»			220	»		
157	ø			189	»			221	»		
158	x			190	»			222	»		
159	f			191	»			223	»		
								224	»		
								225	»		
								226	»		
								227	»		
								228	»		
								229	»		
								230	»		
								231	»		
								232	»		
								233	»		
								234	»		
								235	»		
								236	»		
								237	»		
								238	»		
								239	»		
								240	»		
								241	»		
								242	»		
								243	»		
								244	»		
								245	»		
								246	»		
								247	»		
								248	»		
								249	»		
								250	»		
								251	»		
								252	»		
								253	»		
								254	»		
								255	»		

Memory operands

RISC-V interpreter

RISC-V Interpreter

Input your RISC-V code here:

```
1 lui t0, 0x0
2 lui t1, 0x7ffff
3
4
5
6
7
8
9
10
11
12
13
14
15
16
```

Reset Step Run CPU: 32 Hz

The most recent instructions will be shown here when stepping.

Features

- Reset to load the code, Step one instruction, or Run all instructions
- Set a breakpoint by clicking on the line number (only for Run)
- View registers on the right, memory on the bottom of this page

Supported instructions

- Arithmetic: `ADD`, `ADDI`, `SUB`
- Logical: `AND`, `ANDI`, `OR`, `ORI`, `XOR`, `XORI`
- Shift: `SLT`, `SLTI`, `SLTU`, `SLTIU`
- Memory: `LB`, `SB`, `LBU`, `SBU`, `SW`, `LW`, `SW`
- Jumps: `JAL`, `JALR`
- Branches: `BEQ`, `BNE`, `BLT`, `BLTU`, `BGE`, `BGEU`

RISC-V Reference: [riscv-spec-v2.2.pdf](https://riscv.org/specs/riscv-spec-v2.2.pdf)

Init Value	Register	Decimal	Hex	Binary
0	x0 (zero)	0	0x00000000	0x00000000000000000000000000000000
0	x1 (ra)	0	0x00000000	0x00000000000000000000000000000000
0	x2 (sp)	0	0x00000000	0x00000000000000000000000000000000
0	x3 (gp)	0	0x00000000	0x00000000000000000000000000000000
0	x4 (tp)	0	0x00000000	0x00000000000000000000000000000000
0	x5 (t1)	0	0x00000000	0x00000000000000000000000000000000
0	x6 (t2)	0	0x00000000	0x00000000000000000000000000000000
0	x7 (t3)	0	0x00000000	0x00000000000000000000000000000000
0	x8 (t4)	0	0x00000000	0x00000000000000000000000000000000
0	x9 (t5)	0	0x00000000	0x00000000000000000000000000000000
0	x10 (t6)	0	0x00000000	0x00000000000000000000000000000000
0	x11 (t7)	0	0x00000000	0x00000000000000000000000000000000
0	x12 (t8)	0	0x00000000	0x00000000000000000000000000000000
0	x13 (t9)	0	0x00000000	0x00000000000000000000000000000000
0	x14 (t10)	0	0x00000000	0x00000000000000000000000000000000
0	x15 (t11)	0	0x00000000	0x00000000000000000000000000000000
0	x16 (t12)	0	0x00000000	0x00000000000000000000000000000000
0	x17 (t13)	0	0x00000000	0x00000000000000000000000000000000
0	x18 (t14)	0	0x00000000	0x00000000000000000000000000000000
0	x19 (t15)	0	0x00000000	0x00000000000000000000000000000000
0	x20 (t16)	0	0x00000000	0x00000000000000000000000000000000
0	x21 (t17)	0	0x00000000	0x00000000000000000000000000000000
0	x22 (t18)	0	0x00000000	0x00000000000000000000000000000000
0	x23 (t19)	0	0x00000000	0x00000000000000000000000000000000
0	x24 (t20)	0	0x00000000	0x00000000000000000000000000000000
0	x25 (t21)	0	0x00000000	0x00000000000000000000000000000000
0	x26 (t22)	0	0x00000000	0x00000000000000000000000000000000
0	x27 (t23)	0	0x00000000	0x00000000000000000000000000000000
0	x28 (t24)	0	0x00000000	0x00000000000000000000000000000000
0	x29 (t25)	0	0x00000000	0x00000000000000000000000000000000
0	x30 (t26)	0	0x00000000	0x00000000000000000000000000000000
0	x31 (t27)	0	0x00000000	0x00000000000000000000000000000000

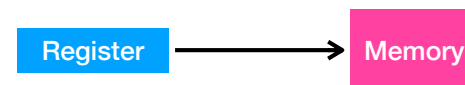
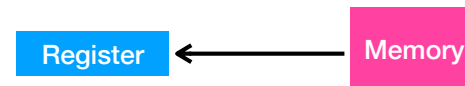
<https://www.cs.cornell.edu/courses/cs3410/2019sp/riscv/interpreter/>

Memory operands

- Main memory used for composite data
 - arrays, structures, dynamic data
- To apply arithmetic operations:
 - **load** values from memory into registers
 - **store** result from register to memory
- Memory is byte addressed
 - each address identifies an 8-bit byte
- RISC-V is **Little Endian**
 - least-significant byte at least address of a word
- RISC-V does not require words to be aligned in memory
 - unlike some other ISAs

Fetching and storing data

- Values can be fetched from memory
 - using a **load** instruction
- Values can be stored in memory
 - using a **store** instruction



Load instruction

- Load doubleword from memory

destination register source address
ld x9, 64(x22)

- destination can be any register
- source address uses a constant (**offset**) added to the register in parentheses (**base**)

Store instruction

- Store doubleword into memory

source register destination address
sd x9, 96(x22)

- source can be any register
- destination address uses a constant (**offset**) added to the register in parentheses (**base**)

Example with memory operand

- C code

```
g = h + A[8];
```

- Compiled RISC-V code

- index 8 requires offset of 64 (8 bytes)

```
# assume g in x21, h in x22  
# assume base address of A in x23  
ld        x21, 64(x23)  
add       x21, x21, x22
```

Example with memory operand

- C code

```
A[12] = h + A[8];
```

- Compiled RISC-V code

- try yourself

```
# assume h in x21  
# assume base address of A in x22  
ld        x9, 64(x22)  
add       x9, x21, x9  
sd        x9, 96(x22)
```

Registers vs memory

- Operating on memory data requires **loads** and **stores**
 - consider memory latency and additional instructions to be executed
- **Registers** are a fast read/write memory right on the CPU that can hold values
- Compiler must use registers for variables as much as possible
 - only spill to memory for less frequently used variables
 - **register optimization** is important!